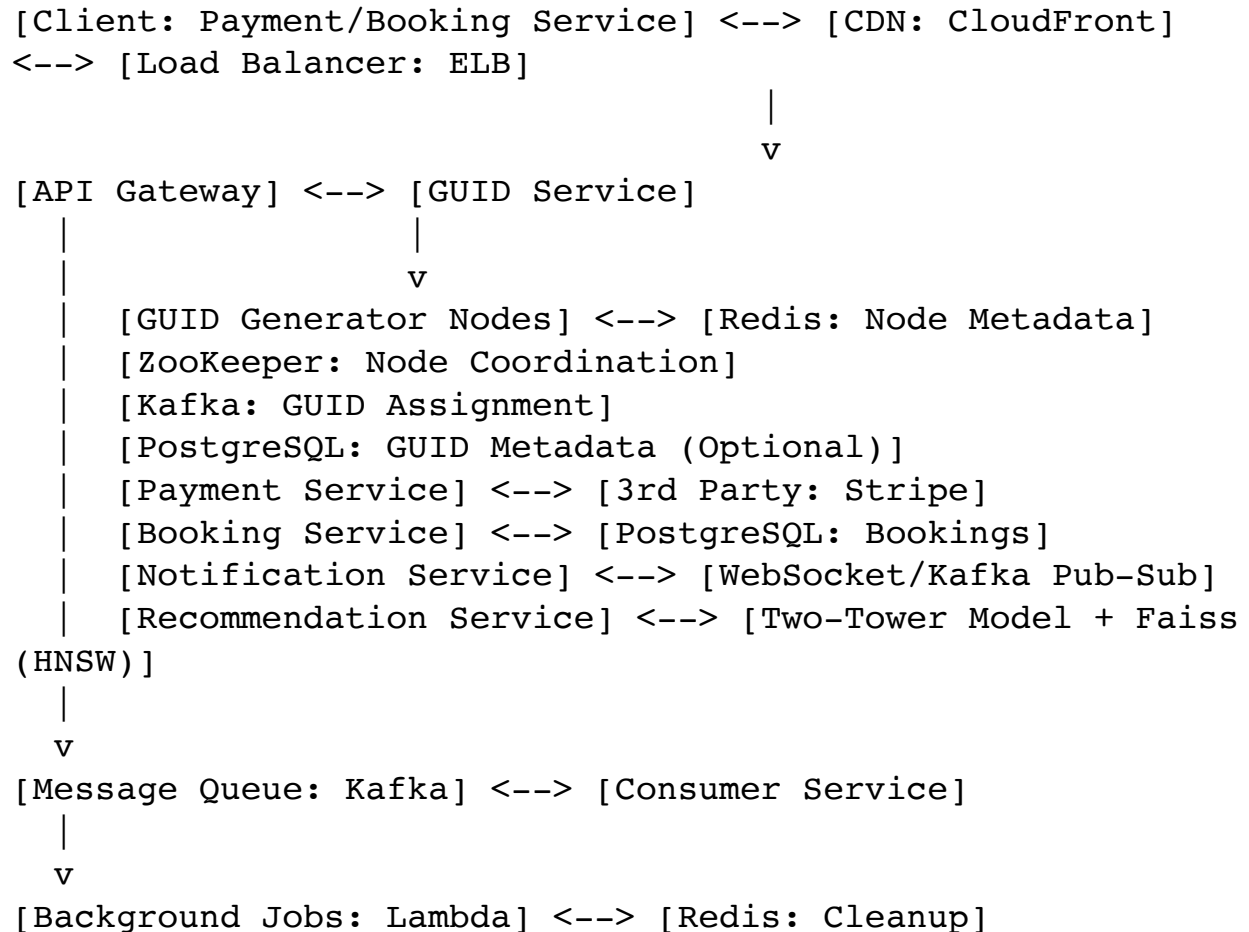


# Distributed GUID System Design

## Architecture Diagram



## Components

### 1. Client

- **Services:** Payment Service, Booking Service, and other microservices in the Buy Ticket Service requesting GUIDs.
- **Interaction:** Call GUID Service via API (e.g., POST /guid/generate) to obtain unique IDs for payments, bookings, etc.
- **Frontend:** Web/mobile apps (React) display GUIDs in booking confirmations.

### 2. CDN

- **Purpose:** Serve static content (e.g., API documentation) for low latency.
- **Tech:** AWS CloudFront, minimal usage as GUID generation is dynamic.

### 3. Load Balancer

- **Purpose:** Distribute 100K QPS across GUID Service instances.
- **Tech:** AWS Elastic Load Balancer (ELB) with Round Robin algorithm.

### 4. API Gateway

- **Purpose:** Route GUID requests, enforce rate limiting (e.g., 1000 requests/sec/client), and handle authentication (JWT optional).
- **Endpoints:**
  - POST /guid/generate?type={uuid|snowflake}: Generate a GUID (UUID or Snowflake-like ID).
  - GET /guid/status/{guid}: Check GUID metadata (e.g., creation time, node).
- **Tech:** AWS API Gateway, rate limiting to prevent abuse.

### 5. GUID Service

- **Purpose:** Generate unique GUIDs across distributed nodes without duplication.
- **Implementation (Snowflake-like ID Generator):**
  - **ID Structure:** 64-bit ID composed of:
    - **Timestamp (41 bits):** Milliseconds since epoch (e.g., 2025-01-01), supports ~69 years.
    - **Node ID (10 bits):** Unique per node (supports 1024 nodes).
    - **Sequence Counter (12 bits):** Per-node counter (supports 4096 IDs/millisecond/node).
    - **Random Bits (1 bit):** Optional for additional entropy.
    - Total:  $41 + 10 + 12 + 1 = 64$  bits.
  - **Generation Logic:**

- Each node generates IDs locally using:
 

```
def generate_snowflake_id(node_id, epoch=2025_01_01_ms):

    timestamp = int(time.time() * 1000) - epoch
    sequence = redis.incr(f"guid:node:{node_id}:sequence") % 4096
    return (timestamp << 23) | (node_id << 12) | sequence
```
- **Node ID Assignment:** ZooKeeper assigns unique node IDs (0-1023) at startup.
- **Sequence Counter:** Redis increments per-node counter (guid:node:{node\_id}:sequence, reset per millisecond).
- **Timestamp:** Synchronized via NTP (Network Time Protocol) to avoid clock drift.
- **Uniqueness Guarantee:**
  - Timestamp ensures IDs are unique over time.
  - Node ID ensures uniqueness across nodes.
  - Sequence counter handles high-frequency requests within the same millisecond.
- **Fault Tolerance:**
  - If a node fails, ZooKeeper reassigns its ID to a new node after a timeout.
  - Clock drift handled by waiting if timestamp decreases (rare, <1ms).
- **Scalability:**
  - $1024 \text{ nodes} \times 4096 \text{ IDs/ms/node} = 4.2\text{M IDs/second}$ , sufficient for 100K QPS.
- **Flow:**
  - Payment Service calls POST /guid/generate → GUID Service node generates Snowflake ID → Returns ID (e.g., 1234567890123456789) in <5ms.

## 6. ZooKeeper

- **Purpose:** Coordinate node IDs and ensure uniqueness across distributed nodes.
- **Implementation:**
  - Maintain a registry of active nodes (/guid/nodes/{node\_id}).
  - Assign unique node IDs (0-1023) to GUID Service instances at startup.
  - Handle node failures by reassigning IDs after a lease timeout (e.g., 30s).
- **Flow:**
  - New GUID node starts → ZooKeeper assigns ID 42 → Node uses ID for Snowflake generation.

## 7. Redis

- **Purpose:** Store sequence counters and handle idempotency for downstream services.
- **Implementation:**
  - **Sequence Counter:** Increment guid:node:{node\_id}:sequence per millisecond (reset via Lua script).  

```
local key = KEYS[1]
```
  - ```
local current_time = redis.call('TIME')[1] * 1000
```
  - ```
local last_time = redis.call('GET', key .. ':last_time') or 0
```
  - ```
if current_time > last_time then
```
  - ```
    redis.call('SET', key .. ':last_time',
```
  - ```
    current_time)
```
  - ```
    redis.call('SET', key, 0)
```
  - ```
end
```
  - ```
return redis.call('INCR', key)
```
  - **Idempotency:** Store processed GUIDs for downstream services (e.g., payments:processed:{payment\_id}, TTL=24h).
- **Flow:**
  - GUID node requests sequence → Redis increments counter → Returns sequence (1ms).

## 8. Kafka

- **Purpose:** Handle async GUID assignment for downstream services (e.g., Payment Service, Booking Service).

- **Implementation:**

- Topic: guid\_assignments for exactly-once delivery of GUIDs to consumers.
- Producer: GUID Service sends GUIDs transactionally:  

```
from kafka import KafkaProducer
```
- ```
producer =  
KafkaProducer(bootstrap_servers=[ 'localhost:9092' ]  
,
```
- ```
enable_idempotence=True,
```
- ```
transactional_id='guid_txn')
```
- ```
producer.init_transactions()  
producer.begin_transaction()  
producer.send('guid_assignments',  
value={'guid':  
'1234567890123456789', 'service': 'payment',  
'booking_id': 'book_1001'})  
producer.commit_transaction()  
Consumer: Payment/Booking Service processes GUIDs transactionally.
```

- **Flow:**

- GUID generated → Kafka enqueues to guid\_assignments → Payment Service consumes (150ms).

## 9. PostgreSQL (Optional)

- **Purpose:** Store GUID metadata for auditing (e.g., creation time, node ID).

- **Schema:**

```
CREATE TABLE GUIDS (  

```

- ```
guid BIGINT PRIMARY KEY,
```
- ```
node_id INT,
```
- ```
created_at DATETIME,
```
- ```
service VARCHAR(50) -- e.g., 'payment', 'booking'
```
- ```
);
```
- **Use Case:** Rarely queried, used for debugging or compliance.

## 10. Downstream Integration

- **Payment Service:**

- Requests GUID for payment\_id → Uses in Stripe payment intent.
- Exactly-once delivery via Kafka transactions.
- **Booking Service:**
  - Consumes GUID for booking\_id → Updates Bookings and Seats tables.
  - Example:
 

```
INSERT INTO Bookings (booking_id, user_id,
event_id, seat_id, status)
```
  - ```
VALUES ('book_1001', 42, 1234, 'A1', 'CONFIRMED')
```
  - ```
ON CONFLICT (booking_id) DO NOTHING;
```
- **Notification Service:**
  - Sends at-most-once notifications (e.g., “Booking ID book\_1001 confirmed”) via WebSocket.
- **Recommendation Service:**
  - Uses Two-Tower Model + Faiss (HNSW) to suggest events if booking fails (~40ms, NDCG@10=0.87).

## 11. Exactly-Once Delivery

- **Approach:** Transactional Messaging (Kafka).
- **Implementation:**
  - GUID Service sends GUIDs to guid\_assignments topic transactionally.
  - Consumer (Payment Service) processes GUIDs and updates PostgreSQL atomically:
 

```
consumer = KafkaConsumer('guid_assignments',
```
  - ```
bootstrap_servers=['localhost:9092'],
```
  - ```
isolation_level='read_committed')
```
  - ```
for msg in consumer:
```
  - ```
    guid_data = msg.value
```
  - ```
    with psycopg2.connect(...) as conn:
```
  - ```
        with conn.cursor() as cur:
```
  - ```
            cur.execute(" " "
```

- `INSERT INTO Payments (payment_id,`  
`booking_id, user_id, amount, status)`
- `VALUES (%s, %s, %s, %s, %s)`
- `ON CONFLICT (payment_id) DO`  
`NOTHING`
- `""", (guid_data['guid'],`  
`guid_data['booking_id'], 42, 200.00, 'PENDING'))`
- `conn.commit()`
- `consumer.commit()`
- **Outcome:** GUIDs assigned exactly once to payments/bookings (150ms).

## 12. Scalability and Concurrency

- **Horizontal Scaling:**
  - Deploy 100+ GUID Service nodes, each with a unique node ID (ZooKeeper).
  - Kafka partitions `guid_assignments` by service or `event_id` for 100K QPS.
- **Concurrency:**
  - Local sequence counters (Redis) avoid contention.
  - ZooKeeper ensures unique node IDs.
- **Traffic Spikes:**
  - Handle 100K QPS with load balancing and Redis caching.
  - Auto-scale GUID Service nodes via Kubernetes.

## 13. Performance Metrics

- **Latency:**
  - GUID generation: ~5ms (timestamp + Redis counter 1ms + processing 4ms).
  - Downstream assignment: ~150ms (Kafka transaction 30ms + PostgreSQL 120ms).
  - Recommendations: ~40ms (Faiss HNSW).
- **QPS:** 4.2M IDs/second (1024 nodes × 4096 IDs/ms), supports 100K QPS.
- **Storage:** 100MB for Redis counters, 6GB/year for PostgreSQL metadata.

- **Uniqueness:** Guaranteed by timestamp + node ID + counter.

## 14. Error Handling

- **Clock Drift:** Wait if timestamp decreases (rare, <1ms).
- **Node Failure:** ZooKeeper reassigns node ID after timeout.
- **Duplicate GUIDs:** Impossible due to unique node IDs and sequence counters.
- **Kafka Failures:** Transactional messaging ensures exactly-once delivery.

## 15. Security

- **Randomness:** Add 1-bit random field to prevent predictable IDs.
- **Authentication:** JWT for API access (optional).
- **Rate Limiting:** API Gateway limits 1000 requests/sec/client to prevent abuse.

## 16. Deployment

- **Cloud:** AWS (EC2, RDS PostgreSQL, ElastiCache Redis, Kafka, ZooKeeper, CloudFront).
- **Monitoring:**
  - Prometheus/Grafana for QPS, latency, and collision rates.
  - CloudWatch for logs, alerts for >10ms latency or node failures.
- **CI/CD:** Docker/Kubernetes for zero-downtime deployments.